

VIT VELLORE

COMPUTER NETWORKS

PROJECT REPORT

**SMART-RUDP: A RELIABLE TRANSPORT PROTOCOL BUILT ON UDP WITH
LEARNING-BASED CONGESTION CONTROL**

*Design, Implementation and Validation in C, with a Multi-Site OSPF Network Design and
Wireshark Analysis*

Submitted by

Harshavardhan Mallela (24BAI0274)

Rohit Arya (24BCE0819)

Shloak Sinha (24BDS0378)

Faculty: Anuradha G

Source code: github.com/Capsturrrr/rudp

8 October 2026

ABSTRACT

The Transmission Control Protocol (TCP) provides reliable, ordered and congestion-aware delivery over an unreliable network, but its mechanisms are usually studied only in theory. This project designs and implements RUDP (Reliable UDP), a custom transport-layer protocol written in C on top of UDP sockets, which reproduces the core mechanisms of TCP: a three-way handshake, sequence numbers with a 16-bit checksum, cumulative acknowledgments, retransmission on timeout and on three duplicate acknowledgments, a Go-Back-N sliding window, and additive-increase/multiplicative-decrease (AIMD) congestion control with slow start and congestion avoidance.

The project is extended in two directions. First, Smart-RUDP replaces the fixed AIMD rule with a tabular Q-learning congestion controller and compares it with AIMD on three emulated paths (5G-like, lossy wireless and satellite-like); it matched but did not beat the rule-based controller on completion time. Second, the protocol is placed in a multi-site network of four OSPF routers, two VLANs, a server LAN and a branch LAN built in Cisco Packet Tracer, and its traffic is analysed in Wireshark with a custom dissector.

The implementation is 673 lines of C in four source files. A test message of 2,442 bytes, sent as 62 packets, was delivered completely at every simulated loss level from 0% to 40%. Over the same range the transfer time rose from 108 ms to 20,236 ms and the final congestion window fell from 18.71 to 2.00 packets, which is the behaviour AIMD is designed to produce. Linux `tc/netem` loss injection and packet capture were used to check the protocol against real operating-system packet loss. The report also records the engineering problems met during the work, with screenshots, and how each was fixed.

Keywords: reliable transport, UDP, TCP, AIMD, congestion control, Q-learning, OSPF, Cisco Packet Tracer, Wireshark.

1. INTRODUCTION

Applications that send data across a network choose between two transport protocols. The User Datagram Protocol (UDP) [8] is simple and fast but gives no guarantee that a datagram arrives, arrives once, or arrives in order. TCP adds all of these guarantees, and also slows its sending rate when the network is congested, but it lives inside the operating-system kernel, which makes its behaviour difficult to observe or change.

The cost of ignoring congestion became clear in October 1986, when the Internet suffered its first congestion collapse and throughput between Lawrence Berkeley Laboratory and UC Berkeley fell from 32 Kbps to 40 bps [1]. TCP congestion control was developed in response. A sender that only retransmits lost data makes such a collapse worse, so reliability and congestion control must be designed together.

This project asks whether TCP's core ideas can be rebuilt from first principles on top of UDP, and whether the result can be shown to work under packet loss. The answer is RUDP, a small user-space protocol written in C and tested on Linux with both a built-in loss simulator and kernel-level network emulation.

1.1 Objectives

- Define a compact packet format with sequence numbers, flags and an integrity check.
- Establish connections with a three-way handshake and close them with a FIN exchange.
- Deliver data reliably and in order using acknowledgments, timeouts and fast retransmit.
- Control the sending rate with AIMD congestion control and make the congestion window visible in logs and graphs.
- Validate the protocol under loss from 0% to 40% and under real packet loss injected by tc/netem.
- Extend the protocol with a Q-learning congestion controller (Smart-RUDP) and compare it honestly with AIMD under emulated 5G, wireless and satellite conditions.
- Build a four-router OSPF network with VLANs, DHCP, an access list and a failover test in Cisco Packet Tracer, and analyse RUDP traffic with Wireshark.
- Record every mistake made during development and how it was corrected.

1.2 Application area

RUDP belongs to user-space reliable transport over UDP, the architecture used by QUIC to carry HTTPS traffic [5][6]. Building on UDP lets a transport evolve without waiting for operating-system or network-operator updates [5]. RUDP is a teaching and testing vehicle, not a production system: because it can be read from end to end, it makes congestion-control behaviour visible, and it provides a platform for studying how a transport reacts when the network misbehaves.

1.3 Purpose and use case

The purpose of the project is a reliable transport that can be read, changed and tested, for the common situation of campus users sending data to a server across a routed network. TCP cannot be modified without changing the operating system, and plain UDP leaves loss, duplication and reordering to the application. A transport built over UDP in user space closes that gap, which is why QUIC and similar protocols are built this way [5][6].

The concrete scenario is a student on the student VLAN uploading data to the server on the server LAN. The access list keeps the student VLAN out of the staff VLAN, and OSPF carries the traffic over the shortest path through R3. If the direct link fails, OSPF reroutes the traffic through R2, which is one hop longer and increases the round-trip time of a running flow. A fixed window rule does not know about this change, while Smart-RUDP observes it through its

RTT-ratio state. The network design of Section 9 is therefore not a separate exercise: it supplies the conditions (routing, access policy and failover) under which the transport is meant to run. The protocol itself was tested on one machine with emulated paths (Sections 5 to 8), because Packet Tracer cannot run the C programs; the report states this wherever it matters. To make the protocol visible to a user, a small two-way messenger (RUDP Chat, src/chat.c) was written on the same packet format: each chat message is one RUDP data segment, delivered in order and acknowledged cumulatively, with simulated loss applied on the sending side. In a test with 30% loss in both directions, all 15 messages of one user and all 12 of the other arrived complete and in order, with the lost segments recovered by timeout and fast retransmit.

2. BACKGROUND AND RELATED WORK

Jacobson and Karels (1988) [1] introduced the packet-conservation principle and added several algorithms to 4BSD TCP, among them round-trip-time variance estimation, exponential retransmit backoff, slow start and congestion avoidance. **Chiu and Jain (1989)** [2] analysed increase and decrease rules and showed that additive increase with multiplicative decrease converges to an efficient and fair state; this is why RUDP grows its window linearly and cuts it by half. **RFC 5681** [3] standardises slow start, congestion avoidance, fast retransmit and fast recovery.

CUBIC [4] replaces linear window growth with a cubic function of the time since the last congestion event, which scales better on fast, long-distance paths. It was not implemented in this project and remains future work. **QUIC** [5][6] shows that reliable, congestion-controlled transport can be built on UDP at Internet scale. RUDP shares its UDP foundation and its user-space design, but has no encryption or stream multiplexing. **NetEm** [7] is the Linux traffic-control module that adds delay, loss, duplication and corruption to packets leaving an interface; it was used here to inject real loss on the loopback interface.

The key researcher for this area is **Van Jacobson**, lead author of [1], whose work underlies every congestion-control mechanism used in RUDP. A fuller literature review accompanies this report.

3. SYSTEM DESIGN

3.1 Overview

RUDP has a client and a server. The client connects to the server on 127.0.0.1, UDP port 8888, performs a handshake, splits a message into 40-byte chunks, sends them under window and congestion control, and closes the connection with a FIN. The server reassembles the chunks in order and acknowledges them. All of this runs over plain UDP sockets.

3.2 Packet format

Every packet starts with a 13-byte header followed by an optional payload of up to 512 bytes. The fields are written into the byte buffer one at a time instead of sending the C structure directly, so compiler padding can never change the layout on the wire.

Table 1: RUDP header layout (13 bytes)

Field	Size	Offset	Purpose
seq_num	4 bytes	0	Sequence number of this packet (counts packets, not bytes)
ack_num	4 bytes	4	Next sequence number the sender of this packet expects
flags	1 byte	8	SYN 0x01, ACK 0x02, FIN 0x04, DATA 0x08; flags may be combined
checksum	2 bytes	9	16-bit integrity check over header fields and payload
payload_len	2 bytes	11	Number of payload bytes that follow

The checksum is a one's-complement sum: all header bytes (except the checksum itself) and payload bytes are added, carries above 16 bits are folded back in, and the result is inverted. A receiver that finds a mismatch drops the packet, and retransmission recovers it.

3.3 Connection management

The client picks a random initial sequence number and sends a SYN. The server replies with a SYN-ACK whose acknowledgment number is the client's number plus one, and the client confirms with an ACK. The client tries up to five times, waiting three seconds for each reply. The server follows a three-state machine: LISTEN, then SYN_RECEIVED after a SYN, then ESTABLISHED after the final ACK. A FIN is acknowledged and returns the server to LISTEN, ready for a new connection. After the fix described in Section 7.1, the server accepts a new SYN in any state.

3.4 Reliable delivery

The client numbers its data packets consecutively. The server accepts only the packet it expects next; it then advances its expected number and returns a cumulative ACK meaning that everything before that number has arrived. An out-of-order or duplicate packet is discarded and the server repeats its last ACK. This is the Go-Back-N strategy.

The client recovers losses in two ways. If three duplicate ACKs arrive, it assumes the packet at the base of the window was lost and resends only that packet (fast retransmit). If no new ACK arrives within 500 ms, it times out and resends every packet from the base of the window to the

last packet sent. A fixed window constant of 4 was used in Phase 4; once congestion control was added in Phase 5, the window became the congestion window and the constant was left unused.

3.5 Congestion control

The client keeps a congestion window (cwnd) and a slow-start threshold (ssthresh), and logs every change to a CSV file. The rules are shown in Table 2.

Table 2: AIMD congestion-control rules

Event	Action
Start	cwnd = 1, ssthresh = 16
New ACK, cwnd below ssthresh (slow start)	cwnd increases by 1 (about doubling every round trip)
New ACK, cwnd at or above ssthresh (congestion avoidance)	cwnd increases by $1/\text{cwnd}$ (about +1 per round trip)
Three duplicate ACKs (fast retransmit)	ssthresh = cwnd / 2, cwnd = ssthresh
Timeout	ssthresh = cwnd / 2, cwnd = 1 (back to slow start)
Limits	cwnd is never below 1 and never above 32

The treatment of fast retransmit is a simplified form of fast recovery: the window is halved and growth resumes in congestion avoidance, without the window inflation described in RFC 5681 [3]. The retransmission timer is a fixed 500 ms rather than an estimate derived from measured round-trip times [1].

3.6 Loss simulation

To test loss on demand, the client accepts a loss percentage as a command-line argument. Before each data packet is sent, including each retransmission, the client draws a random number and silently discards the packet if it falls below the chosen percentage. Handshake and FIN packets and the server's ACKs are never dropped. The random generator is seeded from the clock, so two runs at the same loss level differ.

4. IMPLEMENTATION

4.1 Environment and build

Development used Windows with WSL2 and Ubuntu, the GCC compiler, Make, Git, Valgrind, tc (iproute2), tshark and Wireshark, with Python and matplotlib for graphs. The code is compiled with the warning flags -Wall and -Wextra, debugging symbols and the C11 standard (Figure 1).

```
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ make
mkdir -p bin
gcc -Wall -Wextra -g -std=c11 -o bin/server src/server.c src/common.c
gcc -Wall -Wextra -g -std=c11 -o bin/client src/client.c src/common.c
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$
```

Figure 1: Successful build with make: gcc -Wall -Wextra -g -std=c11 produces bin/server and bin/client with no warnings

4.2 Source files

Table 3: Source files (line counts from the delivered source)

File	Lines	Role
src/common.h	59	Constants, flag bits, packet structure, function prototypes
src/common.c	123	Packing and unpacking, checksum, packet printing
src/client.c	301	Handshake, chunking, window and congestion control, retransmission, logging
src/server.c	190	State machine, in-order delivery, cumulative ACKs, reassembly
Total	673	

4.3 Development phases

Table 4: Implementation phases

Phase	Title	Outcome
0	Environment	WSL2, Ubuntu, toolchain and Git repository set up; build verified
1	Raw UDP echo	Baseline client and server exchanging datagrams
2	Packet format	13-byte header with manual pack/unpack; round-trip and corruption tests; Valgrind clean
3	Handshake	Three-way handshake, state machine and FIN teardown
4	Reliable delivery	Go-Back-N, cumulative ACKs, timeout, fast retransmit; byte-perfect reassembly at 40% simulated loss
5	Congestion control	AIMD with slow start, congestion avoidance and cwnd logging
6	Emulation and capture	Real loss with tc/netem on loopback, captured with tshark and Wireshark
7	Benchmark and graphs	Seven-level loss sweep, CSV output, cwnd and benchmark plots, team handoff

4.4 Version control

The work is stored at github.com/Capsturrrr/rudp. The Phase 7 commit (d9bf3f3, "Phase 7 complete: benchmarking, visualization, resilient server, updated README for team handoff") changed 14 files with 458 insertions and 49 deletions and was pushed to the main branch (Figure 2). No credential is visible in the screenshot; the password prompt was answered at the terminal.

```
[main d9bf3f3] Phase 7 complete: benchmarking, visualization, resilient server, updated README for team handoff
14 files changed, 458 insertions(+), 49 deletions(-)
create mode 100755 benchmark.sh
create mode 100644 cwnd_graph.png
create mode 100644 docs/abstract.pdf
rename Makefile:Zone.Identifier => docs/abstract.pdf:Zone.Identifier (100%)
create mode 100644 docs/benchmark_summary.png
rename README.md:Zone.Identifier => docs/benchmark_summary.png:Zone.Identifier (100%)
create mode 100644 docs/implementation-guide.md
create mode 100644 docs/implementation-guide.md:Zone.Identifier
create mode 100644 logs/benchmark_results.csv:Zone.Identifier
create mode 100644 plot_cwnd.py
Harsh@LAPTOP-IRSGP58S:~/projects/rudp$ git push
Username for 'https://github.com': Capsturrrr
Password for 'https://Capsturrrr@github.com':
Enumerating objects: 22, done.
Counting objects: 100% (22/22), done.
Delta compression using up to 20 threads
Compressing objects: 100% (15/15), done.
Writing objects: 100% (15/15), 151.18 KiB | 25.20 MiB/s, done.
Total 15 (delta 3), reused 0 (delta 0), pack-reused 0 (from 0)
remote: Resolving deltas: 100% (3/3), completed with 3 local objects.
To https://github.com/Capsturrrr/rudp.git
  4f8ae90..d9bf3f3  main -> main
Harsh@LAPTOP-IRSGP58S:~/projects/rudp$
```

Figure 2: Phase 7 commit and push to github.com/Capsturrrr/rudp (4f8ae90..d9bf3f3, main -> main)

5. TEST METHODOLOGY

Three complementary methods were used.

- **Built-in loss simulator.** Drops outgoing client data packets at a chosen percentage (Section 3.6). Used for the benchmark sweep because it is repeatable from a single command.
- **tc/netem on the loopback interface.** The command `tc qdisc add dev lo root netem loss 20%` makes the kernel itself drop 20% of packets. This checks the protocol against loss that the program did not generate. The rule was always checked with `tc qdisc show dev lo` and removed afterwards (Section 7.2).
- **Packet capture.** `tshark` captured the loopback interface inside WSL and the file was opened in Wireshark on Windows, as independent evidence of what crossed the interface.

The benchmark script (`benchmark.sh`) starts the client once for each loss level of 0, 5, 10, 15, 20, 30 and 40 percent with a 30-second limit, measures wall-clock duration with `date` and `awk`, extracts the delivered-packet count and final congestion window from the client output, and writes one CSV row per level. The server must already be running. Each level was run once.

6. RESULTS

6.1 Benchmark sweep

Table 5 and Figures 3 and 4 show the final benchmark. Every level delivered all 62 packets, which together carry the 2,442-byte test message. Transfer time grew from 108 ms with no loss

to 20,236 ms at 40% loss, about 187 times longer, while the final congestion window fell from 18.71 packets to 2.00. The terminal output that produced these numbers is shown in Figure 5.

Table 5: Benchmark results, one run per loss level

Simulated loss	Duration (ms)	Packets delivered	Final cwnd
0%	108	62 / 62	18.71
5%	608	62 / 62	7.30
10%	3,112	62 / 62	3.83
15%	7,118	62 / 62	2.00
20%	7,617	62 / 62	2.90
30%	13,625	62 / 62	2.00
40%	20,236	62 / 62	2.00




Figure 3: Transfer time against simulated loss, from the benchmark in Table 5; all 62 packets were delivered at every level




Figure 4: Final congestion window against simulated loss, from the benchmark in Table 5

```

=== Testing with 0% loss ===
  Duration: 108ms, Packets: 62, Final cwnd: 18.71
=== Testing with 5% loss ===
  Duration: 608ms, Packets: 62, Final cwnd: 7.30
=== Testing with 10% loss ===
  Duration: 3112ms, Packets: 62, Final cwnd: 3.83
=== Testing with 15% loss ===
  Duration: 7118ms, Packets: 62, Final cwnd: 2.00
=== Testing with 20% loss ===
  Duration: 7617ms, Packets: 62, Final cwnd: 2.90
=== Testing with 30% loss ===
  Duration: 13625ms, Packets: 62, Final cwnd: 2.00
=== Testing with 40% loss ===
  Duration: 20236ms, Packets: 62, Final cwnd: 2.00

Benchmark complete. Results in logs/benchmark_results.csv
loss_percent,duration_ms,packets_sent,final_cwnd
0,108,62,18.71
5,608,62,7.30
10,3112,62,3.83
15,7118,62,2.00
20,7617,62,2.90
30,13625,62,2.00
40,20236,62,2.00
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$

```

Figure 5: Terminal output of the final benchmark sweep, with the resulting CSV file

The 20% result (7,617 ms, final cwnd 2.90) is only slightly slower than the 15% result and ends with a larger window. Because each level was run once with a clock-seeded random generator,

such small non-monotonic differences are expected run-to-run variation rather than a protocol effect. The trend across the whole range is clear: more loss means longer transfers and a smaller window.

6.2 Congestion window behaviour

Figure 6 shows the congestion window in a run with no loss. The window starts at 1 and grows by 1 per acknowledgment (slow start) until it reaches the threshold of 16, then grows by about $1/\text{cwnd}$ per acknowledgment (congestion avoidance), ending at 18.71 after 62 packets. The x-axis is the acknowledgment number because the whole transfer takes about a millisecond on the loopback interface.




Figure 6: Congestion window in a run with no loss: slow start up to the threshold of 16, then congestion avoidance

Figure 7 shows separate runs at 0%, 10%, 20% and 40% simulated loss, plotted from the client's own cwnd log. Each timeout (red dot) resets the window to 1 and halves the threshold; each fast retransmit (green triangle) halves the window. Timeouts occur about 500 ms apart, matching the retransmission timer, and after each reset the window grows again in slow start. This sawtooth is the characteristic signature of AIMD. In these particular runs there were 5 timeouts and 2 fast retransmits at 10% loss, 8 and 6 at 20%, and 45 and 2 at 40%. The plots come from one extra run per level made on 8 October 2026 for these graphs, so their event counts differ from run to run and from the benchmark in Table 5.




Figure 7: Congestion window and threshold at 0%, 10%, 20% and 40% simulated loss, with timeouts (red dots) and fast retransmits (green triangles); the 0% panel is plotted against acknowledgment number

6.3 Real loss and packet capture

With tc/netem dropping 20% of loopback packets, the transfer still completed. In the Wireshark/tshark capture, handshake and acknowledgment packets carried 13 bytes (header only) and data packets carried 53 bytes (13-byte header plus 40-byte payload), consistent with the 40-byte chunk size. Retransmission bursts appeared about 500 ms apart, matching the retransmission timeout, and the client logged 12 timeouts in that run. The capture screenshots are not reproduced in this report; the capture file and logs are kept with the project.

7. PROBLEMS ENCOUNTERED AND HOW THEY WERE FIXED

Every real obstacle is recorded here, with the evidence that was collected at the time.

7.1 Server stuck after an abandoned client

Symptom. The first complete benchmark sweep failed identically at every loss level. Each run took between 15.3 and 15.6 seconds and reported no packets and no window size (Figures 8 and 9). The CSV file contained 0 for both fields in every row.

```

=== Testing with 0% loss ===
Duration: 15532ms, Packets: , Final cwnd:
=== Testing with 5% loss ===

```

Figure 8: First symptom: a 15,532 ms run with empty packet-count and window fields

```

^C
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ mv benchmark_fixed.sh benchmark.sh
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ chmod +x benchmark.sh
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ ./benchmark.sh
=== Testing with 0% loss ===
Duration: 15532ms, Packets: , Final cwnd:
=== Testing with 5% loss ===
Duration: 15532ms, Packets: , Final cwnd:
=== Testing with 10% loss ===
Duration: 15433ms, Packets: , Final cwnd:
=== Testing with 15% loss ===
Duration: 15331ms, Packets: , Final cwnd:
=== Testing with 20% loss ===
Duration: 15431ms, Packets: , Final cwnd:
=== Testing with 30% loss ===
Duration: 15632ms, Packets: , Final cwnd:
=== Testing with 40% loss ===
Duration: 15533ms, Packets: , Final cwnd:

Benchmark complete. Results in logs/benchmark_results.csv
loss_percent,duration_ms,packets_sent,final_cwnd
0,15532,0,0
5,15532,0,0
10,15433,0,0
15,15331,0,0
20,15431,0,0
30,15632,0,0
40,15533,0,0

```

Figure 9: Whole sweep failing the same way: every level takes about 15.5 s and the CSV records 0 packets and 0 cwnd

Diagnosis. The command `ss -ulnp | grep 8888` showed a server process still bound to UDP port 8888 (Figure 10), so the server was running but not completing handshakes. The duration matched the client's retry limit: five SYN attempts, each waiting up to three seconds, is 15 seconds, plus start-up overhead. The server only accepted a new SYN while in the LISTEN state. An earlier client had been stopped mid-transfer, which left the server in another state where it ignored every later SYN.

```

40,15533,0,0
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ ss -ulnp | grep 8888
UNCONN 0 0 0.0.0.0:8888 0.0.0.0:* users:((("server",pid=3398,fd=3))
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$

```

Figure 10: Diagnosis: a server process is still bound to UDP port 8888

Fix. The server now accepts a new SYN in any state, resets its sequence tracking and reassembly buffer, and replies with a SYN-ACK. This is the condition if `(pkt.flags & FLAG_SYN)` in `src/server.c`, replacing a test that also required the LISTEN state.

7.2 A leftover netem rule

Symptom. A test run behaved as though it were hung. **Cause.** A `tc/netem` rule from an earlier packet-capture session was still active on the loopback interface (Figure 11) and silently

combined with the loss the benchmark was adding itself. **Fix.** The rule was removed with `sudo tc qdisc del dev lo root netem`, and `tc qdisc show dev lo` then reported only the default noqueue discipline (Figure 12). From then on the qdisc was checked before every test and deleted immediately after every netem test.

```
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ tc qdisc show dev lo
qdisc netem 8001: root refcnt 2 limit 1000 loss 20% seed 4194997339558083968
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$
```

Figure 11: *Leftover rule found: netem with loss 20% still active on lo*

```
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ tc qdisc show dev lo
qdisc netem 8001: root refcnt 2 limit 1000 loss 20% seed 4194997339558083968
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ sudo tc qdisc del dev lo root netem
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ tc qdisc show dev lo
qdisc noqueue 0: root refcnt 2
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$
```

Figure 12: *Rule deleted: only the default noqueue qdisc remains*

7.3 Benchmark script produced wrong numbers

Symptom. One run reported a duration of 5,554,676,559 ms, about 64 days, although the packet count (62) and final window (3.55) were correct (Figure 13). **Cause.** The millisecond timestamp command `date +%s%3N` is not handled correctly by this shell. Separately, an earlier version of the script used `grep -oP`, which is not portable, and left the packet-count and window fields blank. **Fix.** The script now uses `date +%s.%N` with an `awk` calculation for milliseconds, and extracts values with plain `grep` and `sed`. The corrected script produced the results in Section 6.

```
Harsh@LAPTOP-IRSGPS8S:~/projects/rudp$ ./benchmark.sh
=== Testing with 0% loss ===
Duration: 5554676559ms, Packets: 62, Final cwnd: 3.55
=== Testing with 5% loss ===
```

Figure 13: *Timing bug: an impossible duration of 5,554,676,559 ms next to correct packet and window values*

7.4 Cluttered and misplaced project files

Symptom. The project folder in Windows Explorer showed leftover scripts and a stray `server.c` in the top-level directory (Figures 14 to 16), together with many files ending in `:Zone.Identifier`, which Windows adds to files that arrive from elsewhere. **Fix.** `server.c` was moved into `src/`, the extra scripts were deleted (Figure 17), and the `Zone.Identifier` files were removed. The first attempt used the pattern `*.Zone.Identifier` and matched nothing, because the file names contain a colon (Figure 18); the corrected pattern `*:Zone.Identifier` worked (Figure 19), leaving a clean directory of `Makefile`, `README.md`, `benchmark.sh`, `bin`, `cwnd_graph.png`, `docs`, `logs`, `plot_cwnd.py`, `scripts`, `src` and `tests`.

.git	24-09-2026 15:55	File folder	
bin	24-09-2026 15:54	File folder	
docs	28-08-2026 18:46	File folder	
logs	24-09-2026 15:54	File folder	
scripts	28-08-2026 18:46	File folder	
src	31-08-2026 16:26	File folder	
tests	29-08-2026 00:20	File folder	
.gitignore	29-08-2026 00:19	txtfile	1 KB
.gitignore• Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
Makefile	29-08-2026 00:19	File	1 KB
Makefile• Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
README.md	29-08-2026 00:19	MD File	2 KB
README.md• Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
plot_cwnd.py	24-09-2026 20:08	Python Source File	2 KB
benchmark_fixed.sh	24-09-2026 20:23	sh_auto_file	1 KB

Figure 14: Project folder in Windows Explorer, including *benchmark_fixed.sh*

.git	24-09-2026 15:55	File folder	
bin	24-09-2026 15:54	File folder	
docs	24-09-2026 20:39	File folder	
logs	24-09-2026 20:17	File folder	
scripts	28-08-2026 18:46	File folder	
src	31-08-2026 16:26	File folder	
tests	29-08-2026 00:20	File folder	
.gitignore	24-09-2026 20:38	txtfile	1 KB
.gitignore• Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
Makefile	24-09-2026 20:38	File	1 KB
Makefile• Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
README.md	24-09-2026 20:38	MD File	6 KB
README.md• Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
plot_cwnd.py	24-09-2026 20:38	Python Source File	2 KB
benchmark_fixed.sh	24-09-2026 20:23	sh_auto_file	1 KB

Figure 15: Folder listing after file updates; *benchmark_fixed.sh* still present

logs	24-09-2026 20:17	File folder	
scripts	28-08-2026 18:46	File folder	
src	31-08-2026 16:26	File folder	
tests	29-08-2026 00:20	File folder	
.gitignore	24-09-2026 20:38	txtfile	1 KB
.gitignore•Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
Makefile	24-09-2026 20:38	File	1 KB
Makefile•Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
README.md	24-09-2026 20:38	MD File	6 KB
README.md•Zone.Identifier	29-08-2026 00:19	IDENTIFIER File	1 KB
plot_cwnd.py	24-09-2026 20:38	Python Source File	2 KB
benchmark_fixed.sh	24-09-2026 20:23	sh_auto_file	1 KB
benchmark (2).sh	24-09-2026 20:31	sh_auto_file	1 KB
server.c	24-09-2026 20:31	C File	7 KB
benchmark.sh	24-09-2026 20:38	sh_auto_file	1 KB

Figure 16: Stray copies in the top-level directory: *benchmark (2).sh* and *server.c*

```

^C
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ cd ~/projects/rudp
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ mv server.c src/server.c
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ rm -f benchmark_fixed.sh "benchmark (2).sh"
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ ls
Makefile      README.md•Zone.Identifier  benchmark_fixed.sh•Zone.Identifier  cwnd_graph.png  plot_cwnd.py  server_resilient.c•Zone.Identifier
Makefile•Zone.Identifier  benchmark.sh•Zone.Identifier  benchmark_v2.sh•Zone.Identifier  docs            plot_cwnd.py•Zone.Identifier  src
README.md      benchmark.sh               bin                                logs            scripts        tests
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$

```

Figure 17: *server.c* moved into *src/* and the extra benchmark scripts removed

```

Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ rm -f *.Zone.Identifier
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ rm -f server_resilient.c benchmark_v2.sh
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ ls
Makefile      README.md•Zone.Identifier  benchmark_fixed.sh•Zone.Identifier  cwnd_graph.png  plot_cwnd.py  server_resilient.c•Zone.Identifier
Makefile•Zone.Identifier  benchmark.sh•Zone.Identifier  benchmark_v2.sh•Zone.Identifier  docs            plot_cwnd.py•Zone.Identifier  src
README.md      benchmark.sh               bin                                logs            scripts        tests
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$

```

Figure 18: First clean-up attempt: the wildcard **.Zone.Identifier* matched nothing

```

Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ rm -f *:Zone.Identifier
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ rm -f server_resilient.c benchmark_v2.sh
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$ ls
Makefile  README.md  benchmark.sh  bin  cwnd_graph.png  docs  logs  plot_cwnd.py  scripts  src  tests
Harsh@LAPTOP-IRSGPS85:~/projects/rudp$

```

Figure 19: Corrected pattern **:Zone.Identifier* leaves a clean directory

7.5 Wireshark could not see WSL loopback traffic

Capturing on the vEthernet (WSL) interface from the Windows version of Wireshark showed no matching packets, because traffic between two processes inside WSL over 127.0.0.1 never reaches the Windows-visible adapter. The fix was to capture inside WSL with tshark on the lo interface, write a .pcapng file, and copy it to Windows to inspect in Wireshark. The first attempt to write the capture into the home directory failed with permission denied, so it was written to /tmp and copied with elevated rights.

7.6 Wrong assumption about the Windows user name

Copying a capture file to /mnt/c/Users/Harsh/Downloads failed with "No such file or directory" because the Windows account has a different name from the WSL account. Listing /mnt/c/Users showed the real name. The lesson is to confirm a path before using it.

7.7 A credential in a screenshot

A GitHub personal access token appeared in a screenshot shared during setup. A credential that appears in a screenshot or chat should be treated as compromised, so the token was revoked at once and a new one was created and typed only at the terminal prompt. No token appears in any figure in this report.

8. SMART-RUDP: LEARNING-BASED CONGESTION CONTROL

The AIMD policy used so far is a fixed rule. It halves the window on every loss signal, so it cannot tell a random wireless loss from congestion, and it grows by only one packet per round trip once past the threshold, which is slow on a long satellite path. Recent work applies reinforcement learning to this problem [9][10]. This section extends RUDP into Smart-RUDP by replacing the window-update rule with a Q-learning agent [11], and compares the two under emulated network conditions.

8.1 Design

The agent is a tabular Q-learner implemented in a single header (rl_cc.h, about 70 lines). Time is divided into monitor intervals of about one smoothed round-trip time (at least 10 ms). At the end of each interval the agent observes a state, receives a reward for the previous action, updates its table and chooses a new action. Table 6 lists the design.

Table 6: Smart-RUDP agent design

Element	Definition
State (80)	Mean-RTT ratio bucket (4) x retransmission-fraction bucket (4) x window-size bucket (5). The RTT ratio is the interval mean RTT divided by the minimum RTT plus twice the smallest observed RTT deviation, so natural jitter is not mistaken for queueing.
Action (5)	Multiply the window by 0.5, 0.85, 1.0, 1.15 or 1.5. Increases add at least one packet.
Reward	$\sqrt{\text{throughput} / \text{reference}} - 0.5 \times \max(0, \text{RTT ratio} - 1) - 4 \times \text{retransmission fraction}$, in the spirit of the utility functions of PCC [12]. The reference is the bottleneck rate.
Learning	Q(s,a) is updated with learning rate 0.15 and discount 0.9. Exploration probability decays from 0.35 to 0.03 over 40 training episodes. The table starts from a prior that favours growth when the path looks clean and shrinking when it looks congested.

Element	Definition
Safety rule	After a retransmission timeout the window is halved (minimum 2) in Smart-RUDP, and reset to 1 in the AIMD baseline.

Both modes share the same reliability code: Go-Back-N, an adaptive retransmission timeout from the smoothed RTT and its deviation (clamped between 60 and 3,000 ms, doubled after a timeout), and fast retransmit on three duplicate ACKs that resends the whole in-flight window. Only the window policy differs. The baseline therefore differs from the original RUDP client of Section 3 in using the adaptive timeout and in a 128-packet window cap; it keeps slow start to a threshold of 16 followed by additive increase.

8.2 Experimental setup

Real tc/netem needs root and was not available in the build environment, so the path is emulated inside the client: random loss, a bottleneck with a fixed service rate and a tail-drop queue, one-way delay with jitter, and a delayed return path for acknowledgments. Packets still travel through real UDP sockets to the real RUDP server. Three scenarios were used, each trained separately. The loss rates are modest because Go-Back-N degrades sharply at high loss (Section 8.4).

Table 7: Emulated scenarios

Scenario	One-way delay	Jitter	Random loss	Rate	Queue
5G-like	10 ms	3 ms	0.2%	1,500 pkt/s	60 pkt
Lossy wireless	25 ms	10 ms	2%	800 pkt/s	50 pkt
Satellite-like	150 ms	10 ms	0.3%	400 pkt/s	100 pkt

For each scenario the agent was trained for 40 episodes (500 packets per episode, 800 for the satellite case), then its table was frozen and a 1,200-packet transfer was run six times against six runs of the AIMD baseline, using the same starting random seeds.

8.3 Results

Table 8: Mean of six 1,200-packet transfers per scenario (time shown with standard deviation)

Scenario	AIMD time (s)	Smart time (s)	AIMD retx	Smart retx	AIMD RTT (ms)	Smart RTT (ms)
5G-like	1.44 +/- 0.21	2.35 +/- 0.04	62	34	25.1	23.8
Lossy wireless	9.30 +/- 0.89	9.14 +/- 0.79	259	263	59.8	59.5

Scenario	AIMD time (s)	Smart time (s)	AIMD retx	Smart retx	AIMD RTT (ms)	Smart RTT (ms)
Satellite-like	22.0 +/- 2.9	24.3 +/- 2.3	105	88	311.3	309.0




Figure 20: Completion time, retransmitted packets and mean RTT for the rule-based baseline and the frozen Q-learning agent (mean of six runs, error bars are one standard deviation)

Smart-RUDP did not beat the baseline on completion time. It was slower on the 5G-like and satellite-like paths (by about 63% and 10%) and equal within noise on the lossy path. It sent fewer retransmissions on the 5G-like and satellite-like paths and had a slightly lower mean RTT, but the retransmission differences are within one standard deviation across runs and should not be treated as significant. The result is reported as obtained: with a tabular agent, 40 training episodes per scenario and six evaluation runs, learning-based control matched but did not outperform a well-tuned rule on these paths.




Figure 21: Training curves: time to deliver 500 packets in each training episode (dashed line: baseline, scaled to the same length)




Figure 22: Congestion window over time in one evaluation run of each mode and scenario

The training curves in Figure 21 are noisy and do not show a steady improvement; in the 5G-like case later episodes are on average slower than early ones, which suggests that the agent is still exploring and that the training budget is too small for this table. Figure 22 shows that in the 5G-like and satellite-like runs the learned window stays within a narrow band (about 10 to 20 packets) while AIMD climbs to higher peaks and falls back, which is consistent with the learned policy being more conservative and finishing later. In the lossy run the learned window changes rapidly from one interval to the next.

8.4 Development notes and limits

- The first training round was far worse than the baseline (about four times slower on the satellite path). Three changes were made before the reported run: the prior was corrected to encourage growth under light loss, the window is halved rather than reset to one after a timeout, and the throughput term of the reward was square-rooted so that its signal is not drowned by the loss penalty at low utilisation. The reported numbers are from the final configuration and were not selected from several runs of it.
- Go-Back-N limits what any window policy can achieve: with loss of a few percent, most windows contain a loss and the receiver discards everything after it. Selective acknowledgments would help both modes.
- The emulated path keeps packets in order, so reordering was not tested. The TCP and plain-UDP comparisons proposed in the Smart-RUDP abstract were not run.
- A `tc/netem` script (`netem_scenarios.sh`) is supplied for repeating the scenarios with kernel emulation on Linux. It was not run for this report.

9. NETWORK DESIGN AND CISCO PACKET TRACER

To place the transport protocol inside a realistic network, a four-router topology was built in Cisco Packet Tracer (Figures 23 and 24). A campus LAN with two VLANs reaches a server LAN across a routed core of three routers (R1, R2 and R3) connected in a triangle, and a branch LAN hangs off a fourth router (R4) behind R2. All routers run OSPF [13]. The RUDP client would run on a campus PC and the RUDP server on PC2, so every RUDP segment crosses the

routed core as an ordinary UDP datagram to port 8888. In Packet Tracer the routers carry the device labels Router1, Router2, Router3 and Router0 for R1, R2, R3 and R4; the host names configured on them are R1 to R4.




Figure 23: Four-router OSPF network as built: two campus VLANs behind R1, a server LAN behind R3, a branch LAN behind R4, with the interface used on each link




Figure 24: The same network in Cisco Packet Tracer with port labels shown

9.1 Addressing and configuration

Table 9: Addressing plan

Link	Subnet	Notes
R1 Gi0/1 - R2 Gi0/0	10.0.12.0/30	R1 .1, R2 .2
R2 Gi0/1 - R3 Gi0/1	10.0.23.0/30	R2 .1, R3 .2
R1 Gi0/2 - R3 Gi0/2	10.0.13.0/30	R1 .1, R3 .2 (primary path)
R2 Gi0/2 - R4 Gi0/0	10.0.24.0/30	R2 .1, R4 .2
VLAN 10 (students)	192.168.10.0/24	Gateway R1 .1; DHCP pool VLAN10; PC0 192.168.10.10
VLAN 20 (staff)	192.168.20.0/24	Gateway R1 .1; DHCP pool VLAN20; PC1 192.168.20.11
Server LAN	192.168.30.0/24	Gateway R3 .1; PC2 192.168.30.10, PC3 192.168.30.11
Branch LAN	192.168.40.0/24	Gateway R4 .1; DHCP pool BRANCH; PC4 192.168.40.11

Switch1 carries VLANs 10 and 20 to R1 over a trunk, and R1 routes between them with router-on-a-stick sub-interfaces (802.1Q encapsulation). R1 and R4 act as DHCP servers for the campus VLANs and the branch LAN. All four routers run OSPF process 1 in area 0 with router IDs 1.1.1.1 to 4.4.4.4. An extended access list named BLOCK-STUDENT-TO-STAFF on R1 sub-interface Gi0/0.10 denies traffic from the student VLAN to the staff VLAN and permits everything else. Because the replies of connections started from the staff VLAN toward students also have to pass back through that interface, the policy is effectively two-way between those VLANs, while both VLANs can still reach the server. During the build, the cabling of R2 was found by CDP (show cdp neighbors) to differ from the first plan, and the R2 addresses were changed to match the real cables.

9.2 Test plan and observed results

Table 10 lists the checks performed on the built network. The expected values come from the OSPF cost rules (a Gigabit Ethernet link has cost 1), and the right-hand column gives what was observed in Packet Tracer. R1 reaches the server LAN with metric 2 through R3 directly and the branch LAN with metric 3 through R2 (Figure 25). The first ping of every run was lost while ARP resolved, which is normal.


Table 10: Packet Tracer tests and observed results

No.	Check	Expected	Observed
1	R2: show ip ospf neighbor	Three FULL neighbours	1.1.1.1, 3.3.3.3 and 4.4.4.4 reached FULL

No.	Check	Expected	Observed
2	R1: show ip route ospf	192.168.30.0 metric 2, 192.168.40.0 metric 3	As expected: via 10.0.13.2 and via 10.0.12.2
3	PC0: tracert 192.168.30.10	192.168.10.1, 10.0.13.2, server	Exactly these three hops
4	PC4 (DHCP): ping the server	Reachable through R4, R2, R3	4 of 4 replies, TTL 125
5	PC0: ping a VLAN 20 host and the server after the access list	Staff host blocked, server reachable	Staff host: destination unreachable from 192.168.10.1; server still reachable (4 of 4, TTL 126)
6	Shut R1 Gi0/2 during ping -t to the server	A few losses, then recovery via R2	70 of 71 replies; tracert via 10.0.12.2 and 10.0.23.2
7	Restore R1 Gi0/2	Neighbour returns	R1-R3 returned to FULL

Neighbor ID	Pri	State	Dead Time	Address	Interface
3.3.3.3	1	FULL/DR	00:00:36	10.0.13.2	GigabitEthernet0/2
R1#show ip interface brief					
Interface		IP-Address	OK?	Method	Status
GigabitEthernet0/0		unassigned	YES	unset	up
GigabitEthernet0/0.10		192.168.10.1	YES	manual	up
GigabitEthernet0/0.20		192.168.20.1	YES	manual	up
GigabitEthernet0/1		10.0.12.1	YES	manual	up
GigabitEthernet0/2		10.0.13.1	YES	manual	up
Vlan1		unassigned	YES	unset	administratively down
R1#					
01:06:43: %OSPF-5-ADJCHG: Process 1, Nbr 2.2.2.2 on GigabitEthernet0/1 from LOADING to FULL, Loading Done					
01:09:03: %OSPF-5-ADJCHG: Process 1, Nbr 2.2.2.2 on GigabitEthernet0/1 from FULL to DOWN, Neighbor Down: Dead timer expired					
01:09:03: %OSPF-5-ADJCHG: Process 1, Nbr 2.2.2.2 on GigabitEthernet0/1 from FULL to DOWN, Neighbor Down: Interface down or detached					
01:09:03: %OSPF-5-ADJCHG: Process 1, Nbr 3.3.3.3 on GigabitEthernet0/2 from FULL to DOWN, Neighbor Down: Dead timer expired					
01:09:03: %OSPF-5-ADJCHG: Process 1, Nbr 3.3.3.3 on GigabitEthernet0/2 from FULL to DOWN, Neighbor Down: Interface down or detached					
01:09:08: %OSPF-5-ADJCHG: Process 1, Nbr 2.2.2.2 on GigabitEthernet0/1 from LOADING to FULL, Loading Done					
01:09:08: %OSPF-5-ADJCHG: Process 1, Nbr 3.3.3.3 on GigabitEthernet0/2 from LOADING to FULL, Loading Done					
?Bad filename					
%Error parsing filename (Bad file number)					
R1#show ip route ospf					
10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks					
O 10.0.23.0 [110/2] via 10.0.12.2, 00:06:10, GigabitEthernet0/1					
[110/2] via 10.0.13.2, 00:06:10, GigabitEthernet0/2					
O 10.0.24.0 [110/2] via 10.0.12.2, 00:06:10, GigabitEthernet0/1					
O 192.168.30.0 [110/2] via 10.0.13.2, 00:28:57, GigabitEthernet0/2					
O 192.168.40.0 [110/3] via 10.0.12.2, 00:06:00, GigabitEthernet0/1					
R1#					

Figure 25: R1 console: OSPF routes (metric 2 to the server LAN, 3 to the branch LAN), then the Gi0/2 shutdown and restore used for the failover test



The screenshot shows a virtual machine window titled 'PC0' with tabs for 'Physical', 'Config', 'Desktop', 'Programming', and 'Attributes'. The 'Desktop' tab is active, displaying a 'Command Prompt' window. The command prompt shows the results of a series of ping tests to the IP address 192.168.30.10. The first 20 replies are successful with a time of less than 1ms and a TTL of 125. The next two replies show a time of 13ms and 15ms respectively, with a TTL of 125. The final 10 replies are successful with a time of less than 1ms and a TTL of 125. The ping statistics for 192.168.30.10 are displayed, showing 71 packets sent, 70 received, and a 2% loss. The approximate round trip times in milliseconds are shown as Minimum = 0ms, Maximum = 15ms, and Average = 0ms. The user then enters 'Control-C' and '^C' to interrupt the command. Finally, the user enters 'C:\>tracert 192.168.30.10', which shows the tracing route to 192.168.30.10 over a maximum of 30 hops. The route is shown as follows:

```
Tracing route to 192.168.30.10 over a maximum of 30 hops:

  1  0 ms      0 ms      0 ms      192.168.10.1
  2  0 ms      0 ms      0 ms      10.0.12.2
  3  0 ms      0 ms      0 ms      10.0.23.2
  4  0 ms      0 ms      0 ms      192.168.30.10

Trace complete.

C:\>configure terminal
Invalid Command.

C:\>
```

Figure 26: PC0 during the failover test: 70 of 71 replies, then tracert showing the longer path through R2 and R3

```

Reply from 192.168.30.10: bytes=32 time<1ms TTL=125
Reply from 192.168.30.10: bytes=32 time=2ms TTL=125
Reply from 192.168.30.10: bytes=32 time<1ms TTL=125
Reply from 192.168.30.10: bytes=32 time<1ms TTL=125
Reply from 192.168.30.10: bytes=32 time<1ms TTL=125
Reply from 192.168.30.10: bytes=32 time<1ms TTL=125

Ping statistics for 192.168.30.10:
    Packets: Sent = 71, Received = 70, Lost = 1 (2% loss),
    Approximate round trip times in milli-seconds:
        Minimum = 0ms, Maximum = 15ms, Average = 0ms

Control-C
^C
C:\>tracert 192.168.30.10

Tracing route to 192.168.30.10 over a maximum of 30 hops:

  0  0 ms    0 ms    0 ms    192.168.10.1
  1  0 ms    0 ms    0 ms    10.0.12.2
  2  0 ms    0 ms    0 ms    10.0.23.2
  3  0 ms    0 ms    0 ms    192.168.30.10

Trace complete.

C:\>configure terminal
Invalid Command.

C:\>ping 192.168.20.x
Ping request could not find host 192.168.20.x. Please check the name and try again.
C:\>ping 192.168.20.11

Pinging 192.168.20.11 with 32 bytes of data:

Reply from 192.168.10.1: Destination host unreachable.
Reply from 192.168.10.1: Destination host unreachable.

```

Figure 27: *PC0 pinging the staff PC (192.168.20.11): blocked by the access list on R1*

Test 6 is the convergence measurement. When an interface is shut down, the router detects the loss at once and OSPF recomputes its routes within seconds; a failure that the interface does not detect would instead wait for the 40-second OSPF dead interval. Here one echo reply was lost out of 71. The relevance to the transport layer is that the failover path is one hop longer, so the round-trip time of a running RUDP flow increases and the flow must re-learn its window. Smart-RUDP observes exactly this signal through its RTT-ratio state.

Packet Tracer is a simulator: its devices cannot be captured by Wireshark, and the transport tests of Sections 5 to 8 do not run inside it. The Packet Tracer part of the project therefore validates the network design, and the Wireshark part (Section 10) validates the protocol traffic.

10. WIRESHARK PACKET ANALYSIS

Two kinds of capture are used. The first is a real capture of the loopback interface with tshark inside WSL (Section 6.3). The second is a trace exported by Smart-RUDP itself (the --pcap option), which records every data segment and acknowledgment that reaches its destination across the emulated path, as seen at the receiver, in the standard pcap format. The exported file uses the topology addresses (192.168.10.10 to 192.168.30.10, UDP port 8888) so that it reads as a capture on the server LAN of Figure 23. The headers in this file are written by the program

and not by a network interface, and the connection handshake bypasses the emulator and is not included.

A Wireshark dissector written in Lua (`rudp.lua`) decodes the 13-byte RUDP header into sequence number, acknowledgment number, flags (SYN, ACK, FIN, DATA), checksum and payload length. Useful display filters are `smartrudp.flags.data == 1` for data, `smartrudp.flags.ack == 1` for acknowledgments, and `smartrudp.seq == N`, which lists every copy of one packet and so exposes retransmissions.

Figure 28 is a real capture of the client and server on the loopback interface (`rudp_capture_loss.pcapng`, 164 packets) taken with `tshark` while packet loss was applied, and opened in Wireshark with the dissector. The first three packets are the three-way handshake (SYN, SYN-ACK, ACK), and the expanded header of the SYN shows the sequence number, acknowledgment number, flags (0x01), checksum and payload length decoded field by field. In Figure 29, the filter `smartrudp.seq == 1327198699` shows the same data segment twice: once at 0.0004 s and again at 0.501 s. The segment before it was lost, so the receiver kept acknowledging the same number, the sender timed out after 500 ms and resent the whole window (Go-Back-N), which is the retransmission behaviour described in Section 3.4.




Figure 28: Real loopback capture opened with the RUDP dissector: SYN, SYN-ACK and ACK, with the header of the SYN decoded

No.	Time	Source	Destination	Protocol	Length	Info
11	0.000402131	127.0.0.1	127.0.0.1	RUDP	95	[DATA] seq=1327198699 ack=0 len=40
16	0.501092999	127.0.0.1	127.0.0.1	RUDP	95	[DATA] seq=1327198699 ack=0 len=40

Figure 29: Display filter `smartrudp.seq == 1327198699`: the same data segment sent at 0.0004 s and resent at 0.501 s after a timeout




Figure 30: Sequence-against-time view built from the exported traces of a 400-packet transfer over the lossy-wireless path: first copies of data, retransmissions and the cumulative ACK for AIMD (left) and Smart-RUDP (right)

Figure 30 plots the two traces in the style of a Wireshark time-sequence graph. In this run the receiver saw 487 data segments and 487 acknowledgments for the AIMD flow, of which 87 data segments carried a sequence number already seen, and 495 data segments and 495 acknowledgments for the Smart-RUDP flow, of which 95 were repeats. The staircase of repeated points shows Go-Back-N resending a window after each loss. The two flows delivered the same 400 packets; the counts are from a single run, are counted at the receiver (segments lost on the path are not visible), and are not a statistical comparison.

11. LIMITATIONS

- **Go-Back-N.** One lost packet forces retransmission of the whole in-flight window on timeout, which wastes bandwidth compared with Selective Repeat.

- **Fixed timeout in the original client.** The 500 ms retransmission timeout of the benchmarked client is not estimated from measured round-trip time and does not back off; the Smart-RUDP client of Section 8 adds an adaptive timeout with backoff.
- **Simplified fast recovery.** The window is halved on three duplicate ACKs, but window inflation from RFC 5681 is not implemented.
- **Limited test conditions.** All tests ran on one machine over the loopback interface. The simulator drops only client data packets, and delay, jitter and reordering were not tested.
- **Single runs.** Each benchmark level was run once with a clock-seeded random generator, so small differences between levels are not statistically meaningful.
- **No encryption or multiplexing.** RUDP does not provide the security or stream features of QUIC.

12. CONCLUSION AND FUTURE ENHANCEMENTS

RUDP shows that TCP's core guarantees, ordered and complete delivery over a lossy datagram service together with a self-limiting sending rate, can be built and verified from first principles. The final system delivered all 62 packets of a 2,442-byte message at every simulated loss level from 0% to 40%. The congestion window responded as AIMD predicts, shrinking from 18.71 to 2.00 packets as loss rose, and the 500 ms retransmission interval was visible both in the cwnd graph and in the packet capture.

Smart-RUDP then tested whether a learned window policy could improve on the rule. A tabular Q-learning agent matched AIMD on the lossy path and was slower on the 5G-like and satellite-like paths, with somewhat fewer retransmissions; the honest finding is that this simple agent does not beat a well-tuned rule. The network design adds a four-router OSPF network with VLANs, DHCP, an access list and a tested link failover in Cisco Packet Tracer, and the Wireshark work adds a dissector and traces that make the protocol behaviour visible packet by packet.

The project also taught that reliable transport is largely timers and edge cases, that congestion control is a feedback loop in which the sender listens to the network, that independent evidence such as a packet capture matters, and that debugging the environment is as important as debugging the code.

Future enhancements are:

- Selective Repeat or selective acknowledgments, so only lost packets are resent.
- An adaptive retransmission timer based on measured round-trip time, with exponential backoff.
- A CUBIC window-growth variant to compare with the AIMD results reported here.

- Repeating the Smart-RUDP scenarios with kernel tc/netem and on a real multi-machine path, with reordering, and running the TCP and plain-UDP baselines.
- Repeated runs per loss level with averages and confidence intervals.
- A stronger learning agent (more training, richer state, or a deep network as in [9][10]) combined with selective acknowledgments, since the tabular agent here only matched AIMD.

REFERENCES

- [1] V. Jacobson and M. J. Karels, "Congestion Avoidance and Control," in Proc. ACM SIGCOMM '88, 1988, pp. 314-329.
- [2] D.-M. Chiu and R. Jain, "Analysis of the Increase and Decrease Algorithms for Congestion Avoidance in Computer Networks," Computer Networks and ISDN Systems, vol. 17, no. 1, pp. 1-14, 1989.
- [3] M. Allman, V. Paxson and E. Blanton, "TCP Congestion Control," RFC 5681, IETF, Sept. 2009.
- [4] S. Ha, I. Rhee and L. Xu, "CUBIC: A New TCP-Friendly High-Speed TCP Variant," ACM SIGOPS Operating Systems Review, vol. 42, no. 5, pp. 64-74, July 2008.
- [5] A. Langley et al., "The QUIC Transport Protocol: Design and Internet-Scale Deployment," in Proc. ACM SIGCOMM '17, Los Angeles, CA, Aug. 2017.
- [6] J. Iyengar and M. Thomson, Eds., "QUIC: A UDP-Based Multiplexed and Secure Transport," RFC 9000, IETF, May 2021.
- [7] S. Hemminger, "Network Emulation with NetEm," Open Source Development Lab, Apr. 2005, as cited in the tc-netem manual page.
- [8] J. Postel, "User Datagram Protocol," RFC 768, IETF, Aug. 1980.
- [9] K. S. Midhula and P. Arun Raj Kumar, "An Adaptive Congestion Control Protocol for Wireless Networks Using Deep Reinforcement Learning," IEEE Trans. Network and Service Management, vol. 21, no. 2, pp. 2027-2043, 2024.
- [10] Y. Cao, J. Nie, Y. Fan, X. Shao and G. Lei, "DRLFcc: Deep Reinforcement Learning-empowered Congestion Control Mechanism for TCP Fast Recovery in High Loss Wireless Networks," in Proc. IEEE GLOBECOM 2023, pp. 4345-4350, 2023.
- [11] C. J. C. H. Watkins and P. Dayan, "Q-learning," Machine Learning, vol. 8, no. 3-4, pp. 279-292, 1992.
- [12] M. Dong, Q. Li, D. Zarchy, P. B. Godfrey and M. Schapira, "PCC: Re-architecting Congestion Control for Consistent High Performance," in Proc. USENIX NSDI 2015, pp. 395-408.

[13] J. Moy, "OSPF Version 2," RFC 2328, IETF, Apr. 1998.