

Review Cheat Sheet: RUDP and Smart-RUDP

Two pages to read before you go in. 9 Oct 2026.

Why this is a project, not a lab exercise

A lab follows a given specification. We chose a problem (build TCP's guarantees on UDP), designed several cooperating parts (13-byte header, handshake, Go-Back-N, cumulative ACKs, fast retransmit, AIMD, loss simulation), measured them (benchmark sweep, window graphs, Wireshark captures with our own dissector, Smart-RUDP vs AIMD), documented what broke and how we fixed it, and used the result in a routed network design and a working chat between two devices. We also report a negative result: Smart-RUDP did not beat AIMD.

One line: a lab exercise sends a datagram; we built the reliability and congestion control that sit on top of it, measured them, and used them in a real application.

Numbers to remember

Item	Value
Header	13 bytes: seq 4, ack 4, flags 1, checksum 2, length 2
Data packet on the wire	53 bytes (ACK or handshake: 13)
Timeout / fast retransmit	500 ms / 3 duplicate ACKs
Window	slow start to 16, then +1 per RTT, halve on dup ACKs, reset to 1 on timeout, cap 32
Benchmark	62 of 62 packets delivered from 0 to 40% loss
Smart-RUDP	Q-learning, 80 states, 5 actions, 40 training episodes, 6 eval runs per path
Smart-RUDP result	Experiments: +63% time on 5G-like, +10% satellite, tie on lossy. Chat (simulated): faster than AIMD on 5 of 6 held-out paths, equal on 1
Network	4 routers, OSPF area 0, 2 VLANs, DHCP, ACL; server LAN metric 2 via R3, branch metric 3 via R2
Failover	shut R1 Gi0/2: 70 of 71 pings, rerouted via R2
Chat	window 8, 300 ms timeout; terminal and browser; laptop and phone over Tailscale

Hard questions, short answers

Why not TCP? TCP is in the kernel and cannot be changed. Building our own makes every mechanism visible. QUIC and HTTP/3 do the same.

Did Smart-RUDP win? Depends on the setting. Main experiments (5G, lossy, satellite): tied on one path, slower on two, and we report that. Chat, retrained on simulated bursty random-loss paths: faster than AIMD on 5 of 6 held-out paths, equal on 1, because AIMD resets to 1 on every timeout. Both start from window 10; results are simulated; the agent may not shrink the window without a congestion signal (without that rule it lost).

Why Go-Back-N? Simpler to build and explain. It wastes bandwidth under loss; Selective Repeat with SACK is future work.

Tests ran on one machine? Yes, with emulated delay, loss and jitter. They do not prove internet performance and the report says so. The laptop-to-phone chat is a small real-network check.

What does Packet Tracer prove? The network design (OSPF, VLANs, ACL, failover). It cannot run C, so the protocol is validated separately with emulation and Wireshark.

What if the peer disappears? The sender keeps retransmitting on timeout. A connection abort after repeated timeouts is not built; it is the fix we would add.

Does the checksum catch errors? 16-bit one's-complement sum like UDP and TCP; a corrupted packet fails and is retransmitted. It cannot catch every error pattern.

Is the loss real? Benchmark and chat loss is simulated in the program. We also captured a real loopback run with tshark and netem loss.

Why no handshake in the exported trace? The export builds synthetic packets. The real tshark capture shows SYN, SYN-ACK, ACK.

Is the Q-learning really learning? Yes: the learning-curve figure shows it. It is small and tabular with limited training.

Who did what? Agree one answer as a team before you go in.

Say the limits first

Go-Back-N wastes bandwidth. Protocol tests are loopback only. Smart-RUDP did not beat AIMD. The agent is small. **Next:** Selective Repeat with SACK, adaptive timeout, real-network tests.